

Group 5 - Bike Rental Project ADS 505

October 17, 2022

```
[867]: # Import dependences

import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn import preprocessing
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import train_test_split, KFold, cross_val_score
import sklearn.metrics as metrics
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, \
    recall_score, f1_score
from sklearn.neural_network import MLPClassifier, MLPRegressor
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import NearestNeighbors, KNeighborsClassifier, \
    KNeighborsRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.svm import SVR
import math
import operator
from prettytable import PrettyTable
warnings.filterwarnings("ignore")

# mpl.rc_file_defaults()
# plt.rcParams.update(plt.rcParamsDefault)
# plt.rcParams['axes.facecolor'] = 'black'

%matplotlib inline
```

```
[868]: # parse_dates=[0]: We give the function a hint that data in the first column
        ↳ contains dates that need to be parsed.
        # This argument takes a list, so we provide it a list of one element, which is
        ↳ the index of the first column

Seoul_Bike_df = pd.read_csv('/Users/JohnnyBlaze/Website Data Sets/SeoulBikeData.
        ↳ csv', encoding='unicode_escape', parse_dates=[0])
```

```
[869]: Seoul_Bike_df.head()
```

```
[869]:      Date  Rented Bike Count  Hour  Temperature(°C)  Humidity(%)  \
0 2017-01-12          254      0         -5.2           37
1 2017-01-12          204      1         -5.5           38
2 2017-01-12          173      2         -6.0           39
3 2017-01-12          107      3         -6.2           40
4 2017-01-12           78      4         -6.0           36

      Wind speed (m/s)  Visibility (10m)  Dew point temperature(°C)  \
0              2.2           2000          -17.6
1              0.8           2000          -17.6
2              1.0           2000          -17.7
3              0.9           2000          -17.6
4              2.3           2000          -18.6

      Solar Radiation (MJ/m2)  Rainfall(mm)  Snowfall (cm)  Seasons  Holiday  \
0              0.0           0.0           0.0  Winter  No Holiday
1              0.0           0.0           0.0  Winter  No Holiday
2              0.0           0.0           0.0  Winter  No Holiday
3              0.0           0.0           0.0  Winter  No Holiday
4              0.0           0.0           0.0  Winter  No Holiday

      Functioning Day
0              Yes
1              Yes
2              Yes
3              Yes
4              Yes
```

```
[870]: Seoul_Bike_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8760 entries, 0 to 8759
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  8760 non-null  datetime64[ns]
1   Rented Bike Count     8760 non-null  int64
2   Hour                  8760 non-null  int64
```

```

3   Temperature(°C)          8760 non-null   float64
4   Humidity(%)              8760 non-null   int64
5   Wind speed (m/s)         8760 non-null   float64
6   Visibility (10m)         8760 non-null   int64
7   Dew point temperature(°C) 8760 non-null   float64
8   Solar Radiation (MJ/m2)   8760 non-null   float64
9   Rainfall(mm)             8760 non-null   float64
10  Snowfall (cm)            8760 non-null   float64
11  Seasons                  8760 non-null   object
12  Holiday                  8760 non-null   object
13  Functioning Day          8760 non-null   object
dtypes: datetime64[ns](1), float64(6), int64(4), object(3)
memory usage: 958.2+ KB

```

```

[871]: Seoul_Bike_df = Seoul_Bike_df.astype({'Rented Bike Count':'float', 'Hour':
      ↪ 'object'})
      # Seoul_Bike_df.info()

```

```

[872]: # Reformat Column Names
Seoul_Bike_df = Seoul_Bike_df.copy()

Seoul_Bike_df.columns = [d.replace(' ', '_').replace('.', '') for d in
      ↪ Seoul_Bike_df.columns]

Seoul_Bike_df = Seoul_Bike_df.rename(columns={'Wind_speed_(m/s)': 'Wind_speed(m/
      ↪ s)', 'Visibility_(10m)': 'Visibility(10m)',
      'Solar_Radiation_(MJ/m2)':
      ↪ 'Solar_Radiation(MJ/m2)', 'Snowfall_(cm)': 'Snowfall(cm)'})

# Print Column Names
for col in Seoul_Bike_df.columns:
    print(col)

```

```

Date
Rented_Bike_Count
Hour
Temperature(°C)
Humidity(%)
Wind_speed(m/s)
Visibility(10m)
Dew_point_temperature(°C)
Solar_Radiation(MJ/m2)
Rainfall(mm)
Snowfall(cm)
Seasons
Holiday
Functioning_Day

```

```
[873]: # Check for Nulls
Seoul_Bike_df.isnull().sum()
```

```
[873]: Date                                0
Rented_Bike_Count                       0
Hour                                    0
Temperature(°C)                         0
Humidity(%)                             0
Wind_speed(m/s)                         0
Visibility(10m)                         0
Dew_point_temperature(°C)               0
Solar_Radiation(MJ/m2)                  0
Rainfall(mm)                           0
Snowfall(cm)                           0
Seasons                                 0
Holiday                                 0
Functioning_Day                         0
dtype: int64
```

```
[874]: Seoul_Bike_df.describe().transpose().style.
      ↪background_gradient(cmap='brg',axis=None)
```

```
[874]: <pandas.io.formats.style.Styler at 0x7fb0f8429b50>
```

```
[875]: # Count of Unique Values

Seoul_Bike_df.nunique().sort_values(ascending=False)
```

```
[875]: Rented_Bike_Count          2166
Visibility(10m)             1789
Dew_point_temperature(°C)    556
Temperature(°C)             546
Date                        365
Solar_Radiation(MJ/m2)       345
Humidity(%)                  90
Wind_speed(m/s)             65
Rainfall(mm)                61
Snowfall(cm)                51
Hour                        24
Seasons                     4
Holiday                     2
Functioning_Day              2
dtype: int64
```

```
[876]: # Unique Object Dtype Values


print(Seoul_Bike_df.iloc[:, -3:].apply(lambda col: col.unique()))
```

```
Seasons          [Winter, Spring, Summer, Autumn]
Holiday          [No Holiday, Holiday]
Functioning_Day  [Yes, No]
dtype: object
```

```
[877]: # Counts of Holiday
```

```
sns.countplot(data=Seoul_Bike_df, x='Holiday', hue='Holiday')
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0)
plt.show()

print(Seoul_Bike_df['Holiday'].value_counts())
print()
```




```
No Holiday      8328
Holiday         432
Name: Holiday, dtype: int64
```

```
[878]: # Counts of Functioning Day
```

```
sns.countplot(data=Seoul_Bike_df, x='Functioning_Day', hue='Functioning_Day')
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0)
plt.show()

print(Seoul_Bike_df['Functioning_Day'].value_counts())
print()
```




```
Yes      8465
No        295
Name: Functioning_Day, dtype: int64
```

[879]: *# Counts of Seasons*

```
sns.countplot(data=Seoul_Bike_df, x='Seasons', hue='Seasons')
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0)


print(Seoul_Bike_df['Seasons'].value_counts())
print()
plt.show()
```

```
Spring    2208
Summer    2208
Autumn    2184
Winter    2160
Name: Seasons, dtype: int64
```



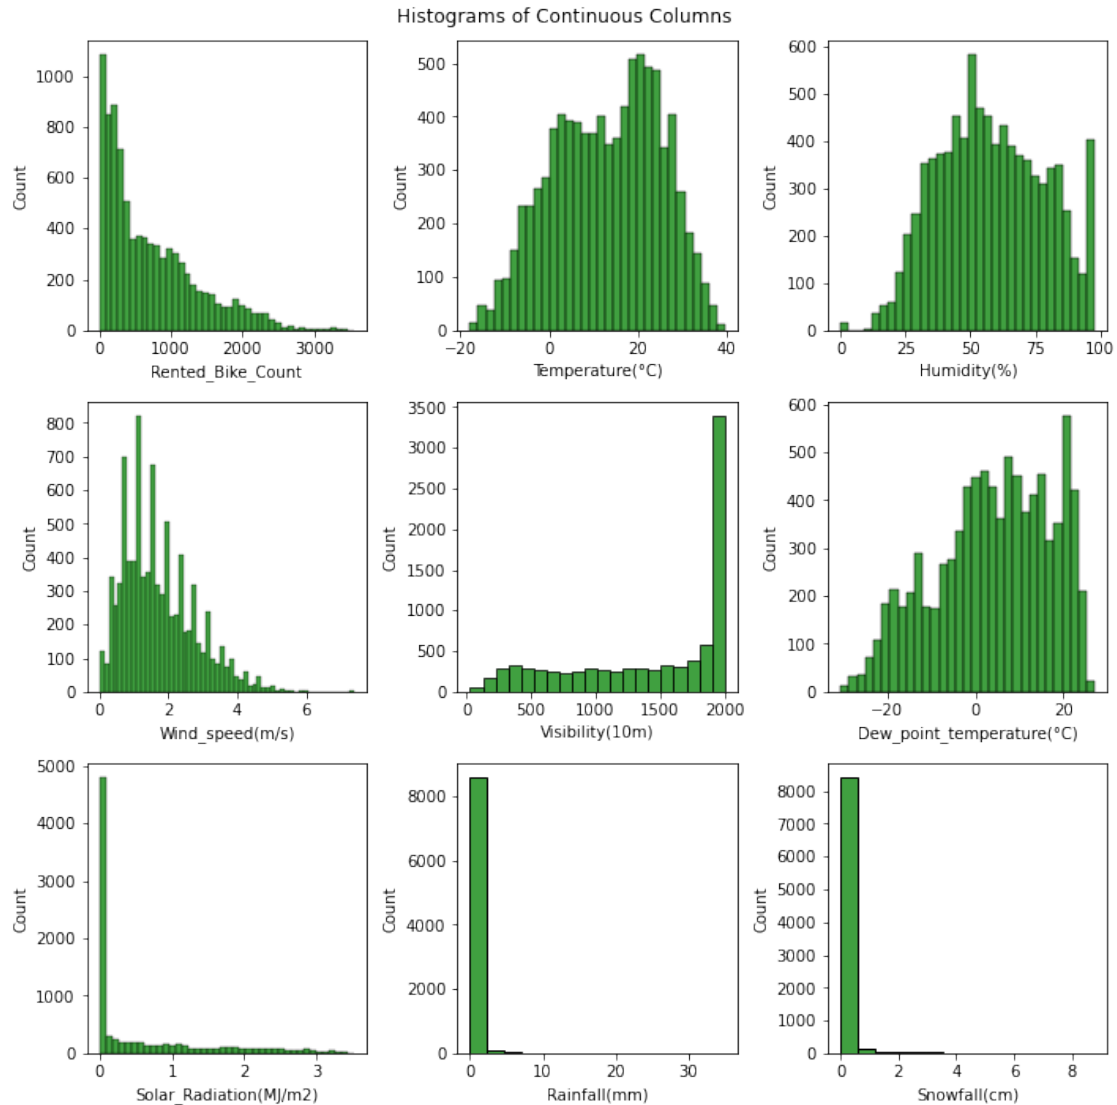
```
[880]: # Bargraphs

plt.figure(figsize=(16, 16))
for i, col in enumerate(Seoul_Bike_df.
    ↳select_dtypes(exclude=['datetime64[ns]', 'object']).columns):
    ax = plt.subplot(3,3, i+1)
    sns.barplot(data=Seoul_Bike_df, x='Hour', y=col, ax=ax, edgecolor='white',
    ↳palette='magma')
plt.suptitle('Bargraphs of Continuous Columns')
plt.tight_layout()
plt.show()
```



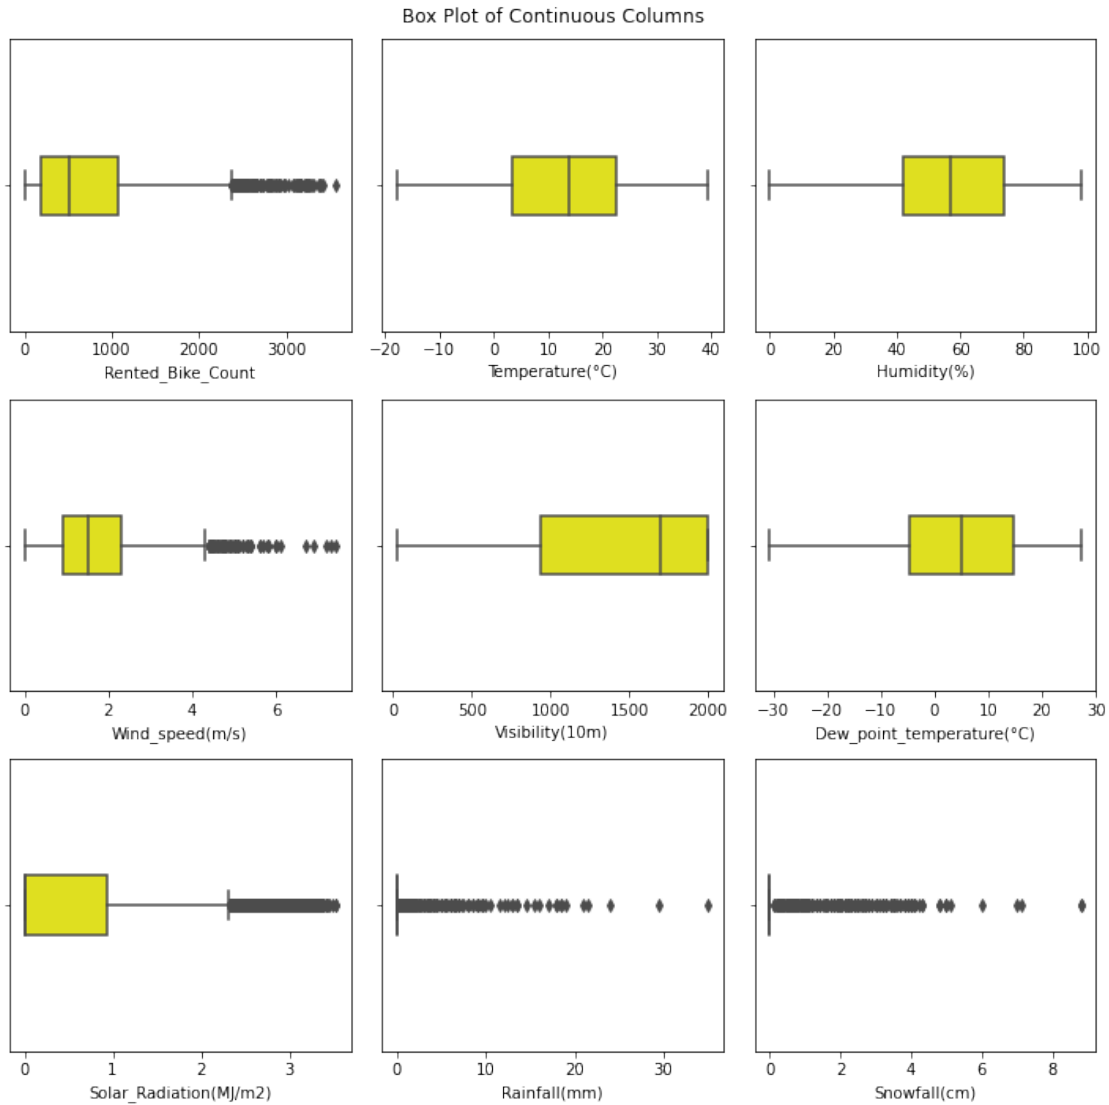
[881]: `# Histograms`

```
plt.figure(figsize=(10, 10))
for i, col in enumerate(Seoul_Bike_df.select_dtypes(include=['float', 'int']).
    columns):
    ax = plt.subplot(3,3, i+1)
    sns.histplot(data=Seoul_Bike_df, x=col, ax=ax, color='green')
plt.suptitle('Histograms of Continuous Columns')
plt.tight_layout()
plt.show()
```

```
[882]: # Box & Whisker

plt.figure(figsize=(10, 10))
for i, col in enumerate(Seoul_Bike_df.select_dtypes(include=['float', 'int']).
    ↪columns):
    ax = plt.subplot(3,3, i+1)
    sns.boxplot(data=Seoul_Bike_df, x=col, ax=ax, color='yellow', width=0.2)
plt.suptitle('Box Plot of Continuous Columns')
plt.tight_layout()
plt.show()
```



[883]: *# Count of Outliers*

```
ContCols = Seoul_Bike_df.select_dtypes(include=['float','int'])

#ContCols.head()


Q1 = ContCols.quantile(0.25)
Q3 = ContCols.quantile(0.75)
IQR = Q3 - Q1

((ContCols < (Q1 - 1.5 * IQR)) | (ContCols > (Q3 + 1.5 * IQR))).sum()
```

```
[883]: Rented_Bike_Count          158
      Temperature(°C)             0
      Humidity(%)                 0
      Wind_speed(m/s)            161
      Visibility(10m)             0
      Dew_point_temperature(°C)   0
      Solar_Radiation(MJ/m2)      641
      Rainfall(mm)                528
      Snowfall(cm)                443
      dtype: int64
```

```
[884]: # Correlation Heatmap

plt.figure(figsize=(8, 8))
heatmap = sns.heatmap(ContCols.corr(method='pearson'), vmin=-1, vmax=1,
    ↪annot=True)
heatmap.set_title('Bike Rental Correlation Heatmap', fontdict={'fontsize':12},
    ↪pad=10);
```



[885]: *# Sort Correlation Values*

```
ContCols[ContCols.columns[:]].corr()['Rented_Bike_Count'][:].
    ↪sort_values(ascending=False)
```

```
[885]: Rented_Bike_Count      1.000000
       Temperature(°C)      0.538558
       Dew_point_temperature(°C) 0.379788
       Solar_Radiation(MJ/m2)  0.261837
       Visibility(10m)        0.199280
       Wind_speed(m/s)       0.121108
       Rainfall(mm)         -0.123074
```

```
Snowfall(cm)          -0.141804
Humidity(%)           -0.199780
Name: Rented_Bike_Count, dtype: float64
```

```
[886]: # Converting Categorical to Dummies

# Hour = pd.get_dummies(Seoul_Bike_df.index.hour, prefix='hour')
Seoul_Bike_df = pd.
↳get_dummies(Seoul_Bike_df, columns=['Holiday', 'Seasons', 'Functioning_Day'], drop_first=True)

# Seoul_Bike_df.head()
```

```
[887]: X = Seoul_Bike_df[['Temperature(°C)', 'Humidity(%)', 'Wind_speed(m/
↳s)', 'Visibility(10m)', 'Dew_point_temperature(°C)',
        'Solar_Radiation(MJ/m2)', 'Rainfall(mm)', 'Snowfall(cm)']]
```

```
[888]: # VIF Function

def _calc_vif(X):
    # Multicollinearity detection
    vif = pd.DataFrame()

    # point here suspicious variables or just all variables
    vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↳shape[1])]
    vif["variables"] = X.columns

    return(vif)
```

```
[889]: # Run VIF

_calc_vif(X).sort_values(by=['VIF'], ascending=False) # High VIF from
↳temperature column and dew point
```

```
[889]:
```

| | VIF | variables |
|---|-----------|---------------------------|
| 0 | 29.075866 | Temperature(°C) |
| 4 | 15.201989 | Dew_point_temperature(°C) |
| 3 | 9.051931 | Visibility(10m) |
| 1 | 5.069743 | Humidity(%) |
| 2 | 4.517664 | Wind_speed(m/s) |
| 5 | 2.821604 | Solar_Radiation(MJ/m2) |
| 7 | 1.118903 | Snowfall(cm) |
| 6 | 1.079919 | Rainfall(mm) |

```
[890]: # Remove Dew Point Column
```

```
X = Seoul_Bike_df[['Temperature(°C)', 'Humidity(%)', 'Wind_speed(m/
↪s)', 'Visibility(10m)',

'Solar_Radiation(MJ/m2)', 'Rainfall(mm)', 'Snowfall(cm)']]
```

```
[891]: # VIF again

_calcul_vif(X).sort_values(by=['VIF'], ascending=False)
```

```
[891]:
```

| | VIF | variables |
|---|----------|------------------------|
| 1 | 4.758651 | Humidity(%) |
| 3 | 4.409448 | Visibility(10m) |
| 2 | 4.079926 | Wind_speed(m/s) |
| 0 | 3.166007 | Temperature(°C) |
| 4 | 2.246238 | Solar_Radiation(MJ/m2) |
| 6 | 1.118901 | Snowfall(cm) |
| 5 | 1.078501 | Rainfall(mm) |

```
[892]: # Define Predictor and Outcome

X = Seoul_Bike_df.iloc[:,2:]
y = Seoul_Bike_df['Rented_Bike_Count']

# X.head()
# X.shape
```

```
[893]: # Split the Data - 75% train, 25% test

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↪random_state=12345)
```

```
[894]: # Scaling

sc = StandardScaler()
X_train_scaled = sc.fit_transform(X_train)
X_test_scaled = sc.transform(X_test)
```

```
[895]: # Linear Regression

lin_reg = LinearRegression()
lin_reg.fit(X_train_scaled,y_train)

# Prediction
y_pred_lin = lin_reg.predict(X_test_scaled)
R2_lin = metrics.r2_score(y_test, y_pred_lin).round(4)
mae_lin = metrics.mean_absolute_error(y_test, y_pred_lin).round(4)
mse_lin = metrics.mean_squared_error(y_test, y_pred_lin).round(4)
```

```


rmse_lin = np.sqrt(mse_lin).round(4)

# Printing the metrics
print('Linear Regression Accuracy: ', lin_reg.score(X_test_scaled, y_test).
      round(4))
print('R2 square:', R2_lin)
print('MAE: ', mae_lin)
print('MSE: ', mse_lin)
print('RMSE: ', rmse_lin)

# Scatterplot
plt.scatter(y_test, y_pred_lin)
plt.plot(y_test, y_test, color = 'red')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Linear Regression: Actual vs Predicted')
plt.show()

```

Linear Regression Accuracy: 0.5397
 R2 square: 0.5397
 MAE: 326.7375
 MSE: 195443.5773
 RMSE: 442.09



```
[896]: # Decision Tree


dt_regressor = DecisionTreeRegressor(random_state = 12345)
dt_regressor.fit(X_train_scaled, y_train)

# Prediction
y_pred_dt = dt_regressor.predict(X_test_scaled)
R2_dt = metrics.r2_score(y_test, y_pred_dt).round(4)
mae_dt = metrics.mean_absolute_error(y_test, y_pred_dt).round(4)
mse_dt = metrics.mean_squared_error(y_test, y_pred_dt).round(4)
rmse_dt = np.sqrt(mse_dt).round(4)

# Printing the metrics
print('Decision Tree Regression Accuracy: ', dt_regressor.score(X_test_scaled,
    ↪y_test).round(4))
print('R2 square:', R2_dt)
print('MAE: ', mae_dt)
print('MSE: ', mse_dt)
print('RMSE: ', rmse_dt)

# Scatterplot
plt.scatter(y_test, y_pred_dt)
plt.plot(y_test, y_test, color = 'red')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Decision Tree: Actual vs Predicted')
plt.show()
```

```
Decision Tree Regression Accuracy: 0.7192
R2 square: 0.7192
MAE: 198.4635
MSE: 119256.1356
RMSE: 345.3348
```

```
[897]: # KNN

knn_9 = KNeighborsRegressor(n_neighbors=9)
KNeighborsRegressor(algorithm='auto', leaf_size=40, metric='minkowski',
                    metric_params=None, n_jobs=-1, n_neighbors=9, p=2,
                    weights='uniform')
knn_9.fit(X_train_scaled, y_train)
# print(knn_9)

# Prediction
y_pred_knn = knn_9.predict(X_test_scaled)
R2_knn = metrics.r2_score(y_test, y_pred_knn).round(4)
mae_knn = metrics.mean_absolute_error(y_test, y_pred_knn).round(4)
mse_knn = metrics.mean_squared_error(y_test, y_pred_knn).round(4)
rmse_knn = np.sqrt(mse_knn).round(4)

# Printing the metrics
print('KNN Regression Accuracy: ', knn_9.score(X_test_scaled, y_test).round(4))
print('R2 square:', R2_knn)
print('MAE: ', mae_knn)
print('MSE: ', mse_knn)
print('RMSE: ', rmse_knn)
```

```

# Scatterplot
plt.scatter(y_test, y_pred_knn)
plt.plot(y_test, y_test, color = 'red')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('KNN: Actual vs Predicted')
plt.show()

```


KNN Regression Accuracy: 0.7747

R2 square: 0.7747

MAE: 201.818

MSE: 95671.1062

RMSE: 309.3075



[898]: *# Random Forest Regression*

```

rf_regressor = RandomForestRegressor(n_estimators = 300 , random_state = 12345)
rf_regressor.fit(X_train_scaled, y_train)

# Prediction
y_pred_rf = rf_regressor.predict(X_test_scaled)
R2_rf = metrics.r2_score(y_test, y_pred_rf).round(4)
mae_rf = metrics.mean_absolute_error(y_test, y_pred_rf).round(4)
mse_rf = metrics.mean_squared_error(y_test, y_pred_rf).round(4)

```

```


rmse_rf = np.sqrt(mse_rf).round(4)

# Printing the metrics
print('Random Forest Regression Accuracy: ', rf_regressor.score(X_test_scaled,
    y_test).round(4))
print('R2 square:', R2_rf)
print('MAE: ', mae_rf)
print('MSE: ', mse_rf)
print('RMSE: ', rmse_rf)

# Scatterplot
plt.scatter(y_test, y_pred_rf)
plt.plot(y_test, y_test, color = 'red')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Random Forest: Actual vs Predicted')
plt.show()

```

Random Forest Regression Accuracy: 0.8642
 R2 square: 0.8642
 MAE: 147.5391
 MSE: 57655.2637
 RMSE: 240.1151



```
[899]: # Neural Network


mlp_reg = MLPRegressor(hidden_layer_sizes=(150,100,50), max_iter = 300,
    ↪activation = 'relu',
                        solver = 'adam', random_state = 12345)
mlp_reg.fit(X_train_scaled, y_train)

# Prediction
y_pred_nn = mlp_reg.predict(X_test_scaled)
R2_nn = metrics.r2_score(y_test, y_pred_nn).round(4)
mae_nn = metrics.mean_absolute_error(y_test, y_pred_nn).round(4)
mse_nn = metrics.mean_squared_error(y_test, y_pred_nn).round(4)
rmse_nn = np.sqrt(mse_nn).round(4)

# Printing the metrics
print('Neural Network Regression Accuracy: ', mlp_reg.score(X_test_scaled,
    ↪y_test).round(4))
print('R2 square:', R2_nn)
print('MAE: ', mae_nn)
print('MSE: ', mse_nn)
print('RMSE: ', rmse_nn)

# Scatterplot
plt.scatter(y_test, y_pred_nn)
plt.plot(y_test, y_test, color = 'red')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Neural Network: Actual vs Predicted')
plt.show()
```

```
Neural Network Regression Accuracy:  0.8548
R2 square: 0.8548
MAE:  158.2314
MSE:  61641.2817
RMSE:  248.2766
```



```
[900]: # SVM

regressor = SVR(kernel='rbf')
regressor.fit(X_train_scaled,y_train)


# Prediction
y_pred_svm = regressor.predict(X_test_scaled)
R2_svm = metrics.r2_score(y_test, y_pred_svm).round(4)
mae_svm = metrics.mean_absolute_error(y_test, y_pred_svm).round(4)
mse_svm = metrics.mean_squared_error(y_test, y_pred_svm).round(4)
rmse_svm = np.sqrt(mse_svm).round(4)

# Printing the metrics
print('Support Vector Regression Accuracy: ', regressor.score(X_test_scaled,
↪y_test).round(4))
print('R2 square:', R2_svm)
print('MAE: ', mae_svm)
print('MSE: ', mse_svm)
print('RMSE: ', rmse_svm)

# Scatterplot
plt.scatter(y_test, y_pred_svm)
plt.plot(y_test, y_test, color = 'red')
```

```
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('SVM: Actual vs Predicted')
plt.show()
```

Support Vector Regression Accuracy: 0.3373
 R2 square: 0.3373
 MAE: 353.3543
 MSE: 281388.0445
 RMSE: 530.4602



[901]: *# Table Results*

```
Table = PrettyTable(["Model", "R-Squared", "MAE", "MSE", "RMSE"])
Table.add_row(["Linear Regression", R2_lin, mae_lin, mse_lin, rmse_lin])
Table.add_row(["Decision Tree", R2_dt, mae_dt, mse_dt, rmse_dt])
Table.add_row(["KNN", R2_knn, mae_knn, mse_knn, rmse_knn])
Table.add_row(["Random Forest", R2_rf, mae_rf, mse_rf, rmse_rf])
Table.add_row(["Neural Network", R2_nn, mae_nn, mse_nn, rmse_nn])
Table.add_row(["SVM", R2_svm, mae_svm, mse_svm, rmse_svm])
print("Models Performance Sorted by R-Squared Values")
Table.sortby = "R-Squared"
print(Table)
```

Models Performance Sorted by R-Squared Values

| Model | R-Squared | MAE | MSE | RMSE |
|-------------------|-----------|----------|-------------|----------|
| SVM | 0.3373 | 353.3543 | 281388.0445 | 530.4602 |
| Linear Regression | 0.5397 | 326.7375 | 195443.5773 | 442.09 |
| Decision Tree | 0.7192 | 198.4635 | 119256.1356 | 345.3348 |
| KNN | 0.7747 | 201.818 | 95671.1062 | 309.3075 |
| Neural Network | 0.8548 | 158.2314 | 61641.2817 | 248.2766 |
| Random Forest | 0.8642 | 147.5391 | 57655.2637 | 240.1151 |